



SAPIENZA
UNIVERSITÀ DI ROMA

Hierarchical Task Network Planning

AI 16-17 Section 1, Planning - 3

Thanks to Francesco Riccio,
Guglielmo Gemignani

Hierarchical Task Networks HTN

In real problems, search at the level of primitive actions can be very inefficient.

Often, the actions make very small changes wrt the goal.

- Some steps of the plan may remain abstract until execution.

Examples:

- planning a holiday,
- a mobile robot cleaning the table in a home after dinner.

HTN principles

Basic Idea: create a **hierarchical decomposition** such that:

- The operators of classical planning are preserved and a more abstract concept of action is introduced:

High Level Actions (HLA)

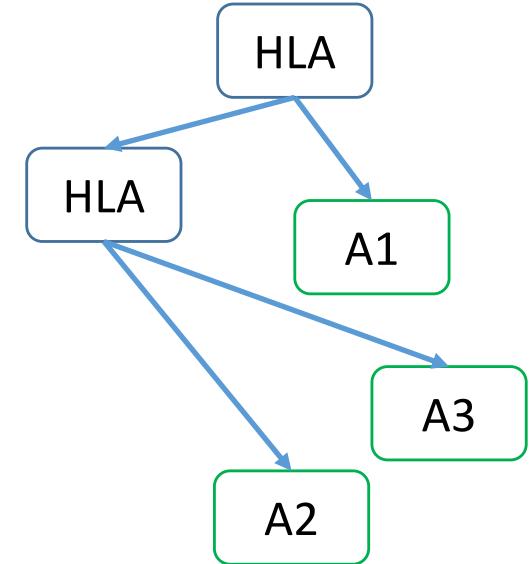
- Planning proceeds by decomposing **nonprimitive** tasks recursively into smaller subtasks, until **primitive** tasks are obtained.

HTN Action Descriptions

Two types of actions:

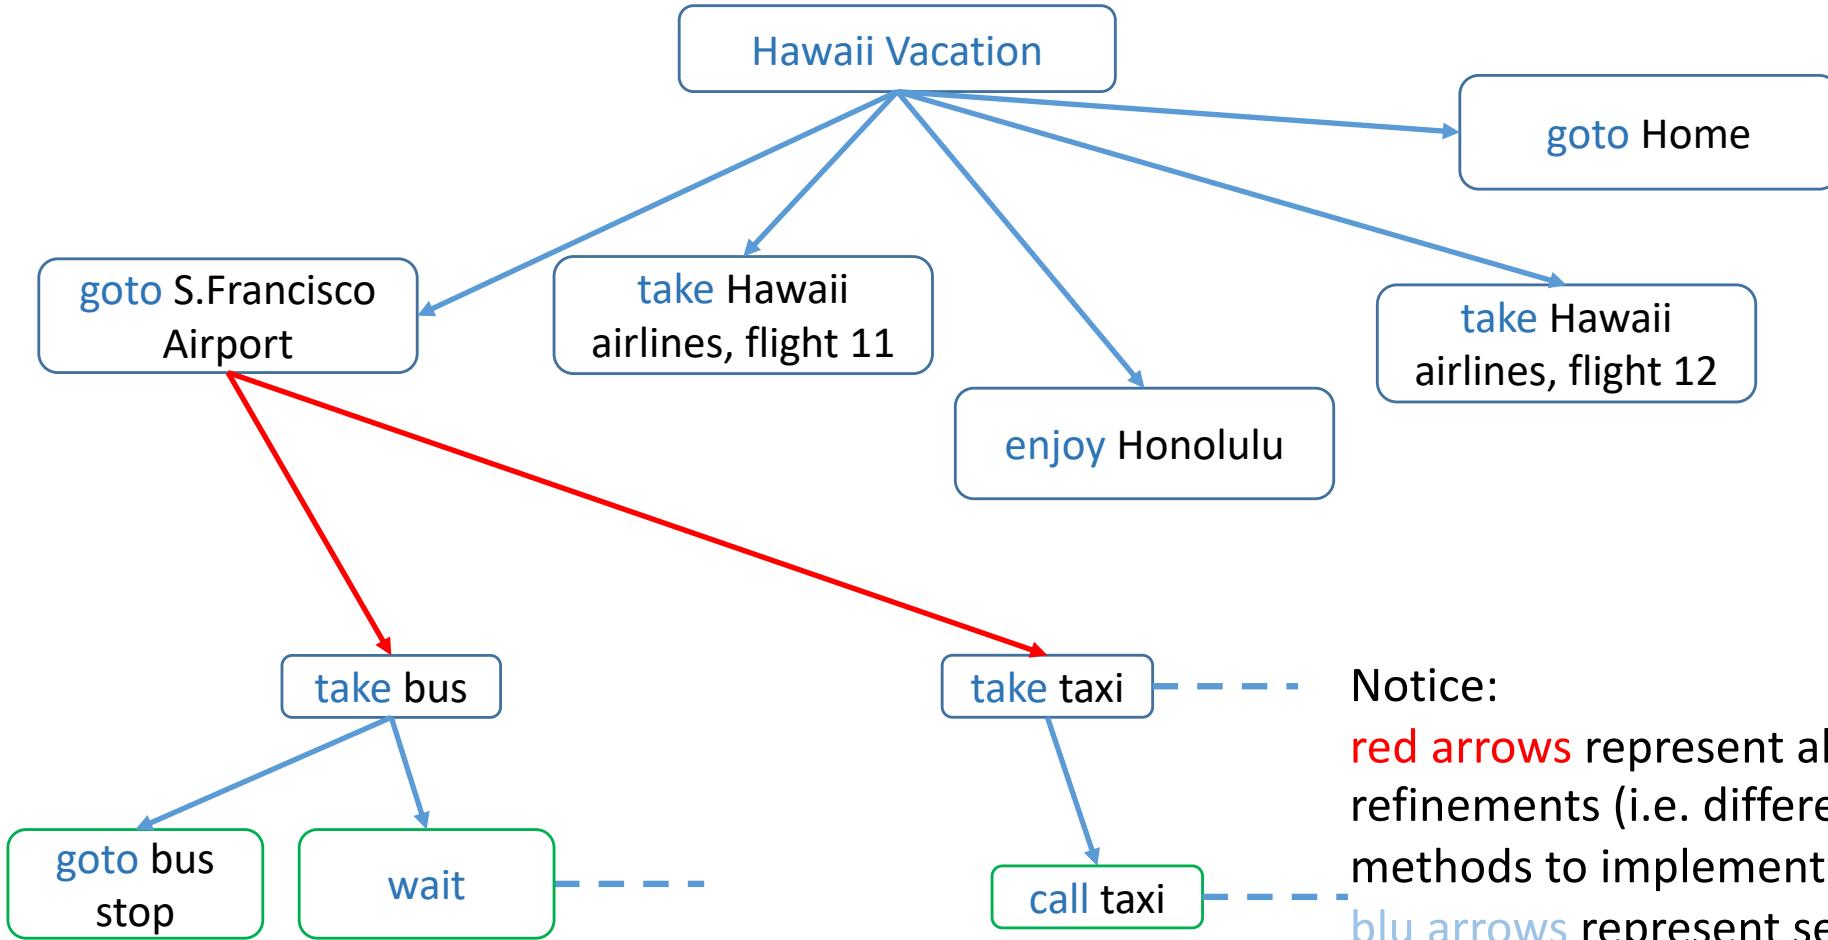
- Primitive actions **A** that (as in classical planning) can be executed by the agent
- High-Level Actions **HLA** that cannot be executed by the agent, but have one or more **refinements** into *sequences of actions*.
(HLA and their refinements embody a plan)

These **sequences of actions** can contain both high-level and primitive actions



HTN Example

The solution is to create a hierarchical decomposition

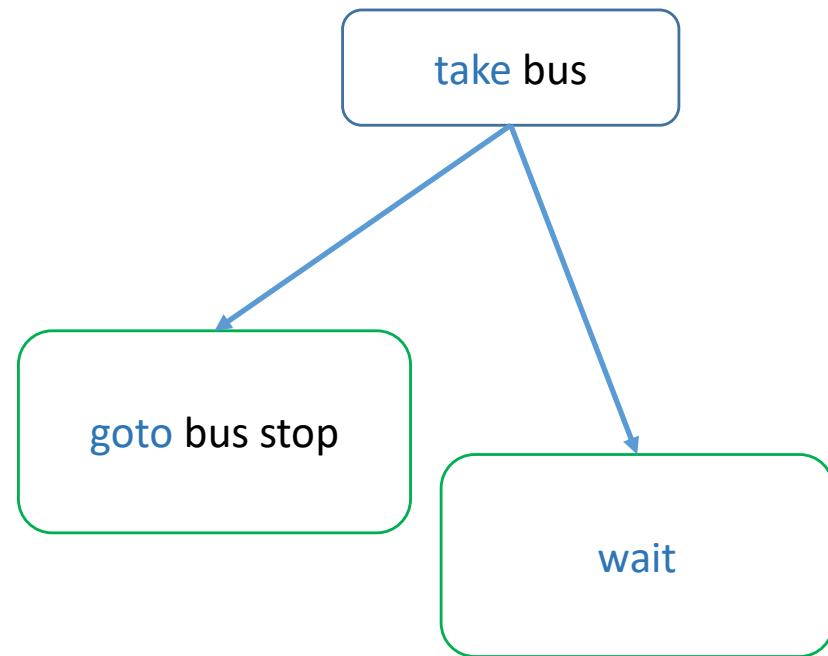


Notice:
red arrows represent alternative refinements (i.e. different methods to implement the HLA)
blue arrows represent sequences of HLA within a single refinement

High Level Actions (HLA)

HLA refinements containing **only primitives** are called **implementations** of the HLA.

An implementation of an **High level plan HLP** is the **concatenation** of the implementations of each HLA in sequence.



Reaching Goals

Given the precondition-effect definitions of each primitive, one can determine whether any given implementation of a HLP achieves the goal!

An High-level plan achieves the goal from a given state if at least one of its implementations does.

Search for plans

Base case:

If an HLA has exactly one implementation, then we can compute the preconditions-effects of the HLA exactly as if it were a primitive action.

What If an HLA has multiple different implementations?

- Search among the implementations
- Search for abstract plans
(Reasoning on HLAs)

A

HLA

Searching for primitive solutions (one implementation for HLA)

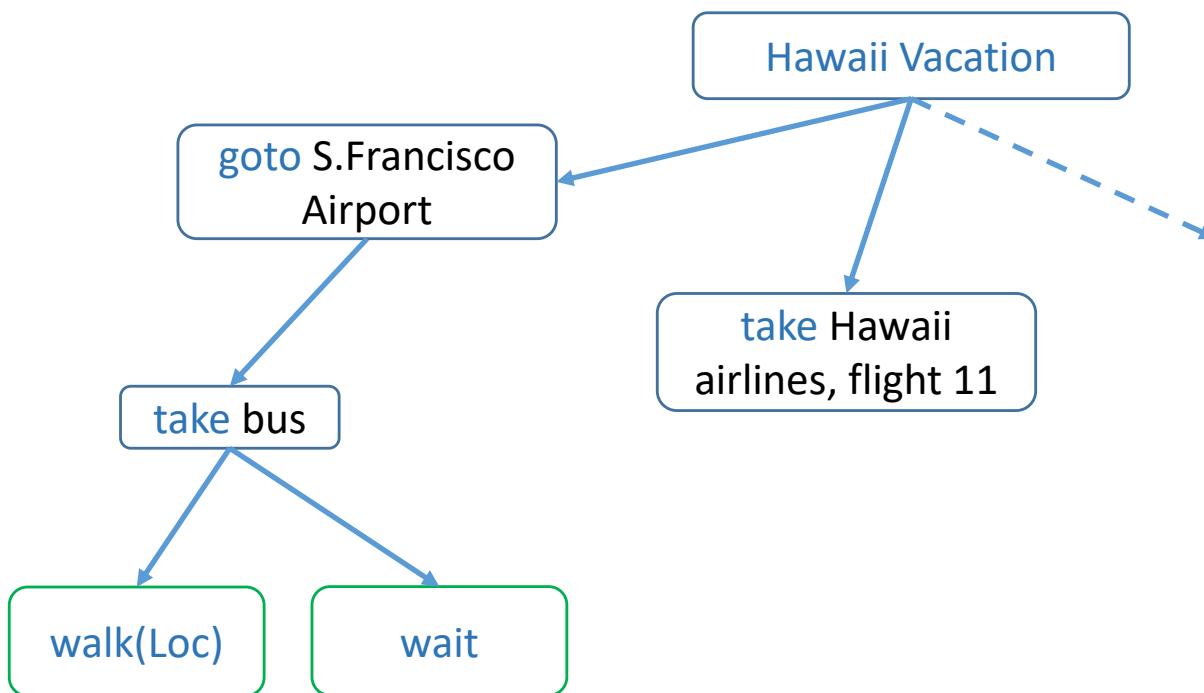
A

The hierarchical search refines HLA from the initial abstract specification to a sequence of primitives to determine whether a plan is workable.

- Repeatedly choose an HLA in the current plan and replace it with one of its refinements until goal is reached.
- The goal check can be performed when the plan is refined to a sequence of primitive actions.
- Complexity $b^d \rightarrow r^{(d-1/k-1)}$,
b=number of actions, d=length of the solution,
r = number of refinements, k = length of refinements

Searching for primitive solutions

A



HPlanning algorithm

A breadth-first implementation
hierarchical forward planning search

function HIERARCHICAL-SEARCH(*problem, hierarchy*) **returns** a solution, or failure

frontier \leftarrow a FIFO queue with [Act] as the only element

loop do

if EMPTY?(*frontier*) **then return** failure

plan \leftarrow POP(*frontier*) /* chooses the shallowest plan in *frontier* */

hla \leftarrow the first HLA in *plan*, or *null* if none

prefix,suffix \leftarrow the action subsequences before and after *hla* in *plan*

outcome \leftarrow RESULT(*problem.INITIAL-STATE, prefix*)

if *hla* is null **then** /* so *plan* is primitive and *outcome* is its result */

if *outcome* satisfies *problem.GOAL* **then return** *plan*

else for each *sequence* **in** REFINEMENTS(*hla, outcome, hierarchy*) **do**

frontier \leftarrow INSERT(APPEND(*prefix, sequence, suffix*), *frontier*)

Searching for abstract plans

HLA

To really take advantage of the power of hierarchical decomposition we need to guarantee that:

abstract plans achieve the goal.

i.e.: The planner should be able to determine that, **eventually**, the HLAs **[Drive(Home, Airport), Fly(Airport, Honolulu)]** get the agent to the goal **without looking for their refinements**.

NOTE: If an High-level plan achieves the goal, working in a small search space of HLAs, then we can let the planner refine each step of the plan.

Downward refinement

HLA

=> Condition for success: a HLP that achieves the goal, if **at least one of its implementation does.**

Downward refinement property for HLAs

Approach: define precondition-effect descriptions of the HLAs as in the case of primitives.

How can we model the effects of a high-level action with multiple refinements?

Conservative approach

HLA

Include **ONLY** the positive effects that are achieved by **every** implementation of the HLA and the negative effects of **any** implementation.

Ex: Without Cash, can't get to the airport ...

- Demonic non-determinism
- Angelic non-determinism

Reachable Sets

HLA

REACH(s,h) of an **HLA h** from **state s**, is the set of states reachable by any of the HLA's implementations.

The reachable set of a **sequence of HLAs** is the **union of all the reachable sets** obtained by applying the HLA n-th in each state in the reachable set of the (n-1)-th HLA, i.e.

$$REACH(s_n, [h_0, \dots, h_n]) = \bigcup_{i=1}^n REACH(s_i, [h_i, \dots, h_0])$$

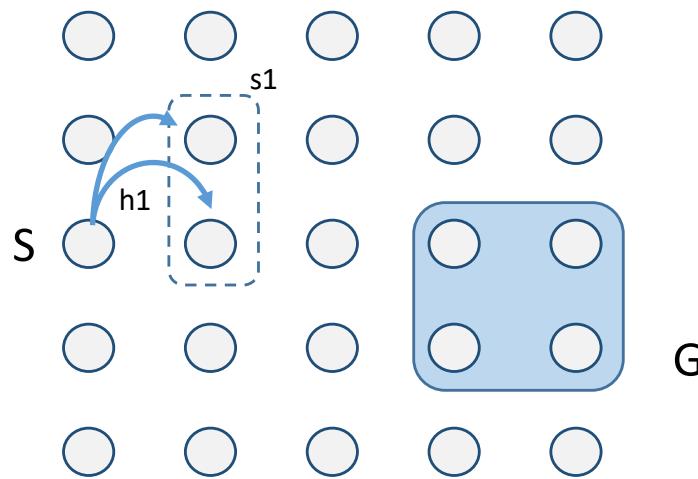
$s_i \in REACH(s_{i-1}, [h_{i-1}, \dots, h_0])$

Reachable sets: example

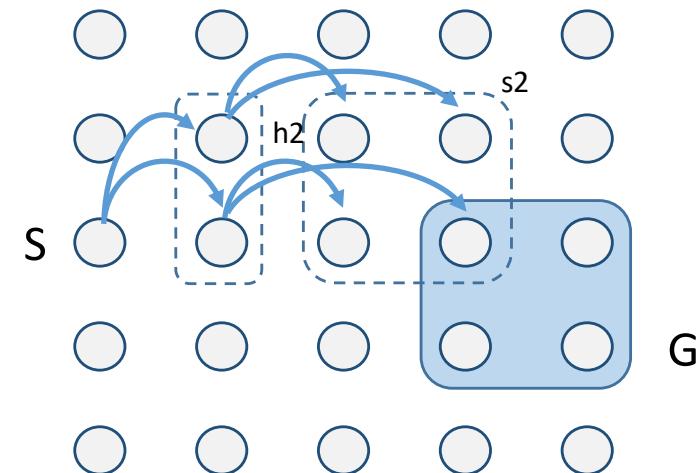
HLA

Given a sequence of HLAs, it achieves the goal if its
reachable set intersects the set of possible goal states.

NB: If it does not intersect, then there is no plan.



REACH(s_1, h_1)



REACH(s_2, h_2)

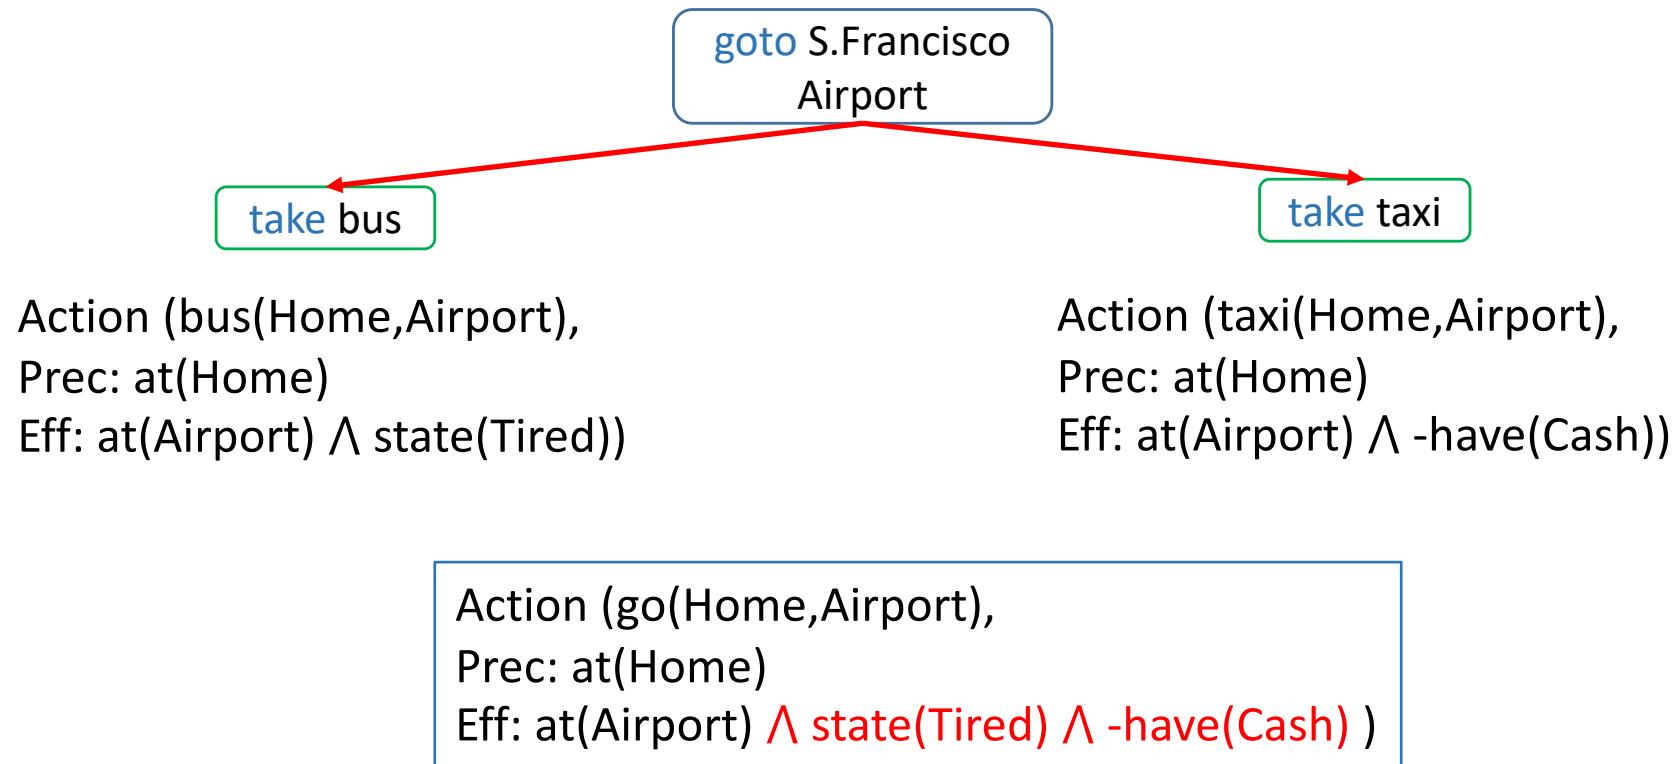
Searching for abstract solutions

HLA

Does the HLA reach the goal?

1. Search among high-level plans;
2. If one reachable set intersects the goal:
Commit to that abstract plan
(knowing that it works)
return yes;
3. Else refining the plan further.
if refinement returns yes, then return yes;
else return no;

Preconditions and effects of HLA



Both will **never** be true

Representing preconditions and effects

HLA

Primitive actions can **add** or **delete** variables or leave them **unchanged**.
An HLA can do more, it can **control** the value of a variable.

+ the **add** symbol
- the **delete** symbol

± the **add-or-delete** symbol
~ the **possibly** symbol

+~ at(airport) means that **possibly adds** at(airport)

Action (bus(Home,Airport),
Prec: at(Home)
Eff: at(Airport) \wedge state(Tired))

Action (taxi(Home,Airport),
Prec: at(Home)
Eff: at(Airport) \wedge -have(Cash))

Action (go(Home,Airport),
Prec: at(Home)
Eff: at(Airport) \wedge -~ have(Cash) \wedge +~ state(Tired))

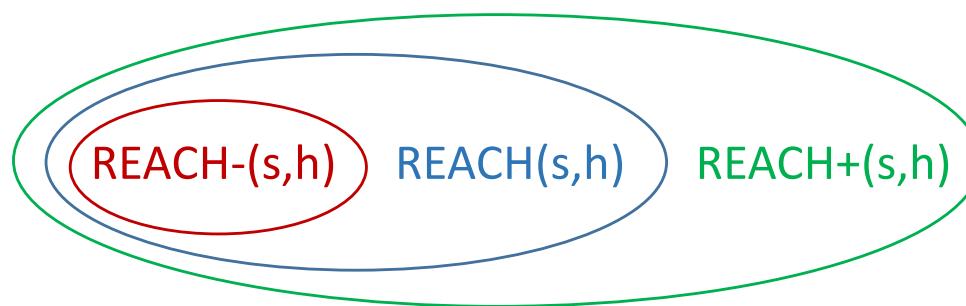
Approximated representation

With infinitely many refinements possible, effects should be approximated

Action (go(Home,Airport),
Prec: at(Home)
Eff: at(Airport), $\neg \sim \text{have}(\text{Cash}) \wedge +\sim \text{state}(\text{Tired})$)

The $\text{REACH}(s,h)$ needs to be redefined with approximate effects:

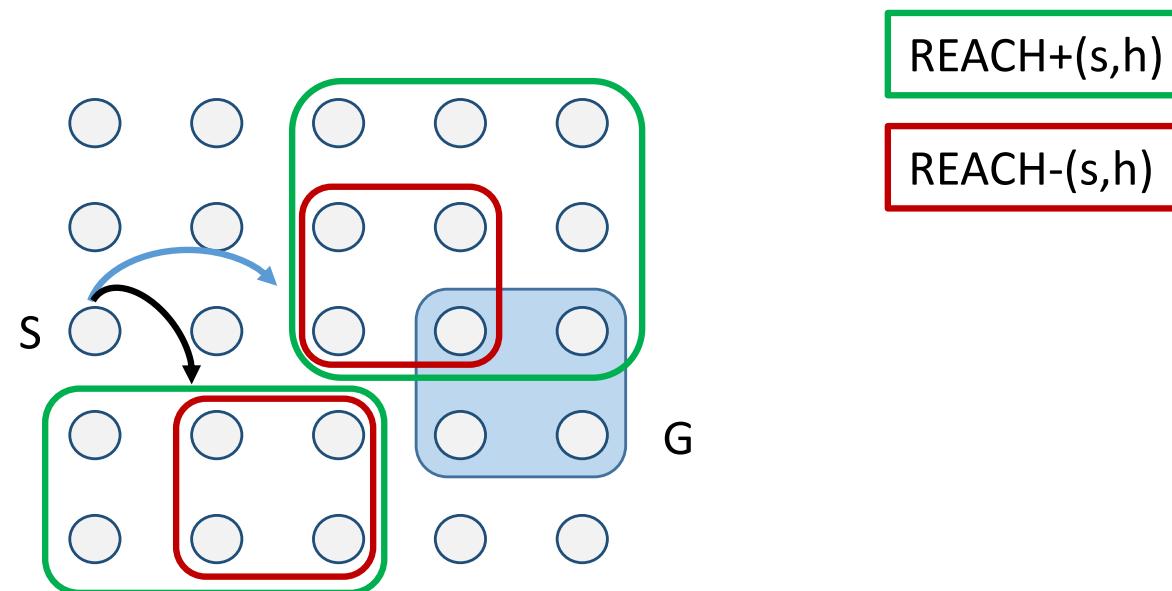
- Optimistic description, overstate the reachable set $\text{REACH+}(s,h)$
Possible effects as definite!
- Pessimistic description, understate the reachable set $\text{REACH-}(s,h)$
Possible effects do not happen!



Plan existence

If the **optimistic** reachable set **does not intersect** the goal
=> the plan does not work

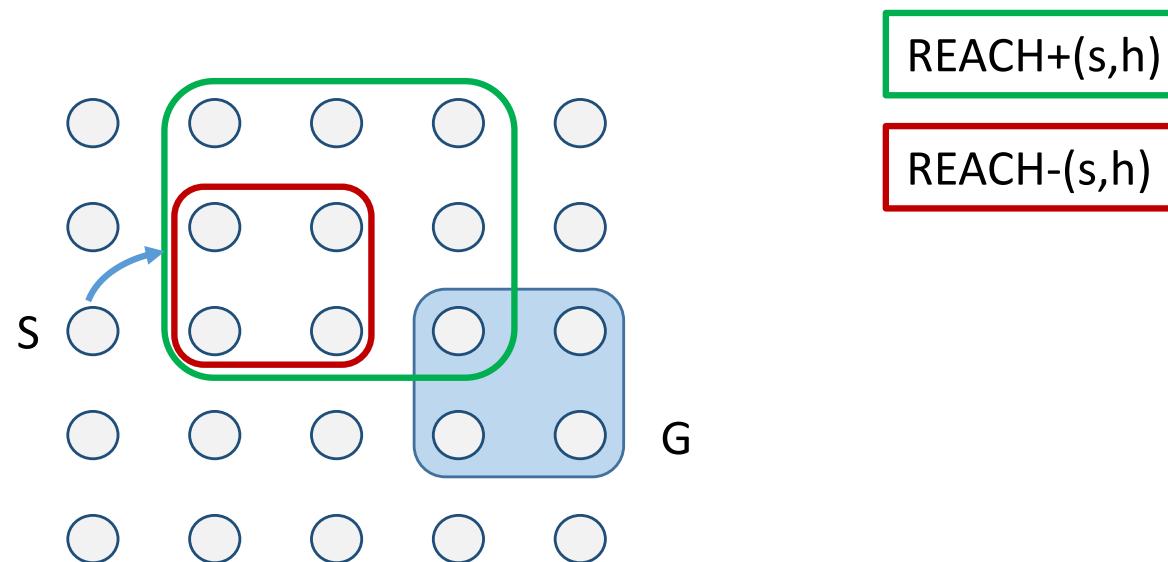
If the **pessimistic** reachable set **intersects** the goal
=> the plan works



Not always possible to determine existence!

If the **optimistic** reachable set **does intersect** the goal
=> the plan might work

If the **pessimistic** reachable set does not **intersect** the goal
=> the plan might work



HPlanning for HLPlans

Basis hierarchical refinement search + approximate effects.

function ANGELIC-SEARCH(*problem*, *HTN*, *initialPlan*) **returns** a solution, or failure

frontier \leftarrow a FIFO queue with *InitialPlan* as the only element

loop do

if EMPTY?(*frontier*) **then return** failure

plan \leftarrow POP(*frontier*) /* chooses the shallowest node in *frontier* */

if REACH⁺ (*problem.INITIAL-STATE*, *plan*) intersects *problem.GOAL* **then**

if *plan* is primitive **then return** *plan*

guaranteed \leftarrow REACH⁻ (*problem.INITIAL-STATE*, *plan*) \cap *problem.GOAL*

if *guaranteed* $\neq \{ \}$ and MAKING-PROGRESS(*plan*, *initialPlan*) **then**

finalState \leftarrow any element of *guaranteed*

return DECOMPOSE(*HTN*, *problem.INITIAL-STATE*, *plan*, *finalState*)

hla \leftarrow some HLA in *plan*

prefix, *suffix* \leftarrow the action subsequences before and after *hla* in *plan*

for each *sequence* in REFINEMENTS(*hla*, *outcome*, *HTN*) **do**

frontier \leftarrow INSERT(APPEND(*prefix*, *sequence*, *suffix*), *frontier*)

HLP Decomposition

```
function DECOMPOSE( $HTN, s_0, plan, s_f$ ) returns a solution
   $solution \leftarrow$  an empty plan
  while  $plan$  is not empty do
     $action \leftarrow$  REMOVE-LAST( $plan$ )
     $s_i \leftarrow$  a state in REACH $^-(s_0, plan)$  such that  $s_f \in$  REACH $^-(s_i, action)$ 
     $problem \leftarrow$  a problem with INITIAL-STATE= $s_i$  and GOAL= $s_f$ 
     $solution \leftarrow$  APPEND(ANGELIC-SEARCH( $problem, HTN, action$ ),  $solution$ )
     $s_f \leftarrow s_i$ 
  return  $solution$ 
```

MAKING-PROGRESS checks that the planner is not in an infinite loop recursively applying the same refinement.

Visit Grandpa

Primitives:

Action (`takeBus`(From, To, Money),
Prec: has(Money) \wedge at(From)
Eff: -has(Money) \wedge -at(From) \wedge
at(To))

Action (`walk`(From, To),
Prec: at(From)
Eff: -at(From) \wedge at(To))

Action (`talk()`,
Eff: state(happy))

Action (`getIcecream`(Money),
Prec: has(Money)
Eff: state(happy)) \wedge -has(Money))



Initial state: has(money1) \wedge
has(money2) \wedge at(home)

Goal state: state(happy)

Visit Grandpa



HLAs:

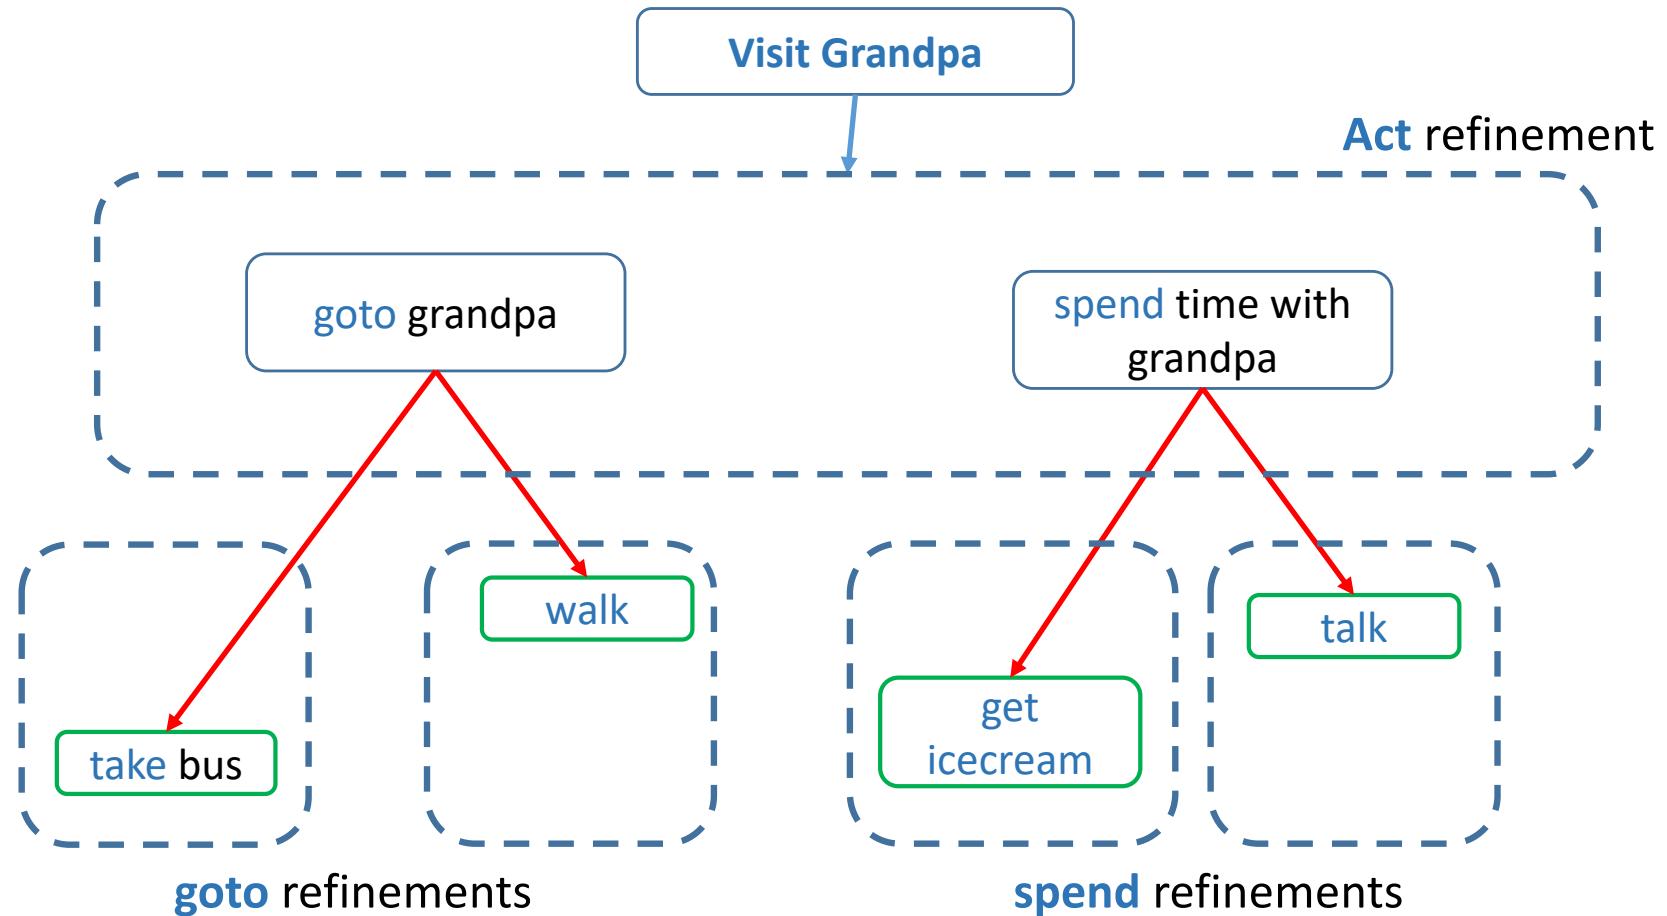
HLA (**goto**(From,To,Money),
Prec: at(From)
Eff: -at(From) \wedge at(To) \wedge -~has(money)
Refinement: [takeBus(),walk()])

Initial state: has(money1) \wedge
has(money2) \wedge at(home)

Goal state: state(happy)

HLA(**spendtime**(Money),
Prec: at(granpa)
Eff: state(happy) \wedge -~has(Money)
Refinement: [getIcecream(),talk()])

Visit Grandpa



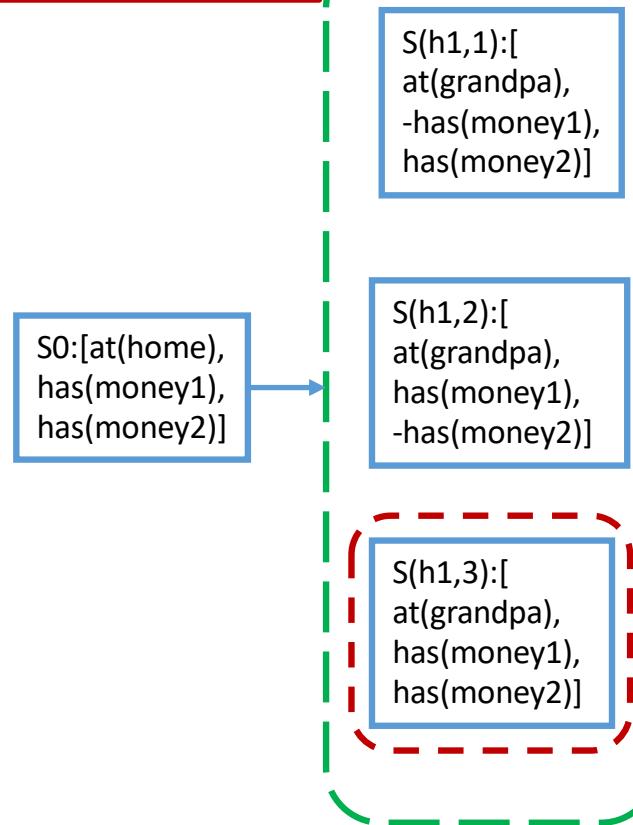
Visit Grandpa

REACH(s,h) \pm

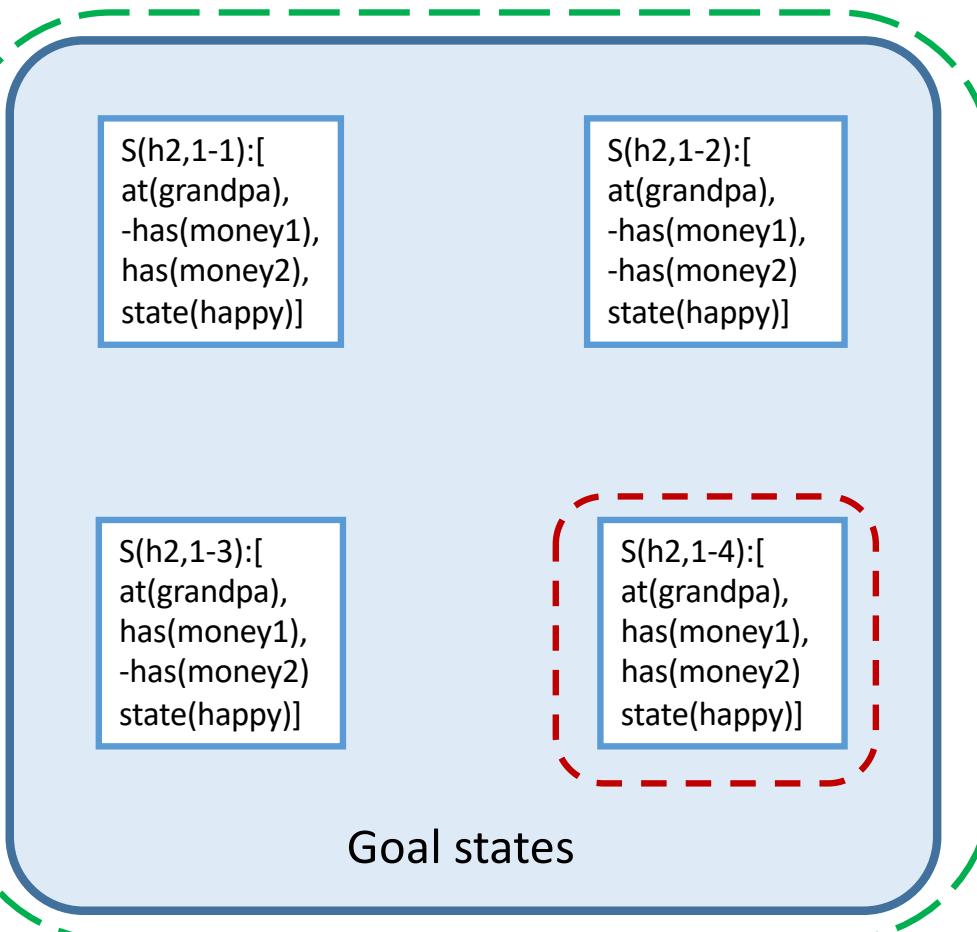
Initial state: has(money1) \wedge
 has(money2) \wedge at(home)
 Goal state: state(happy)



Goal in REACH-
 A PLAN exists



REACH(s0,h1)



REACH(s1,[h1,h2])

REACH+
 REACH-

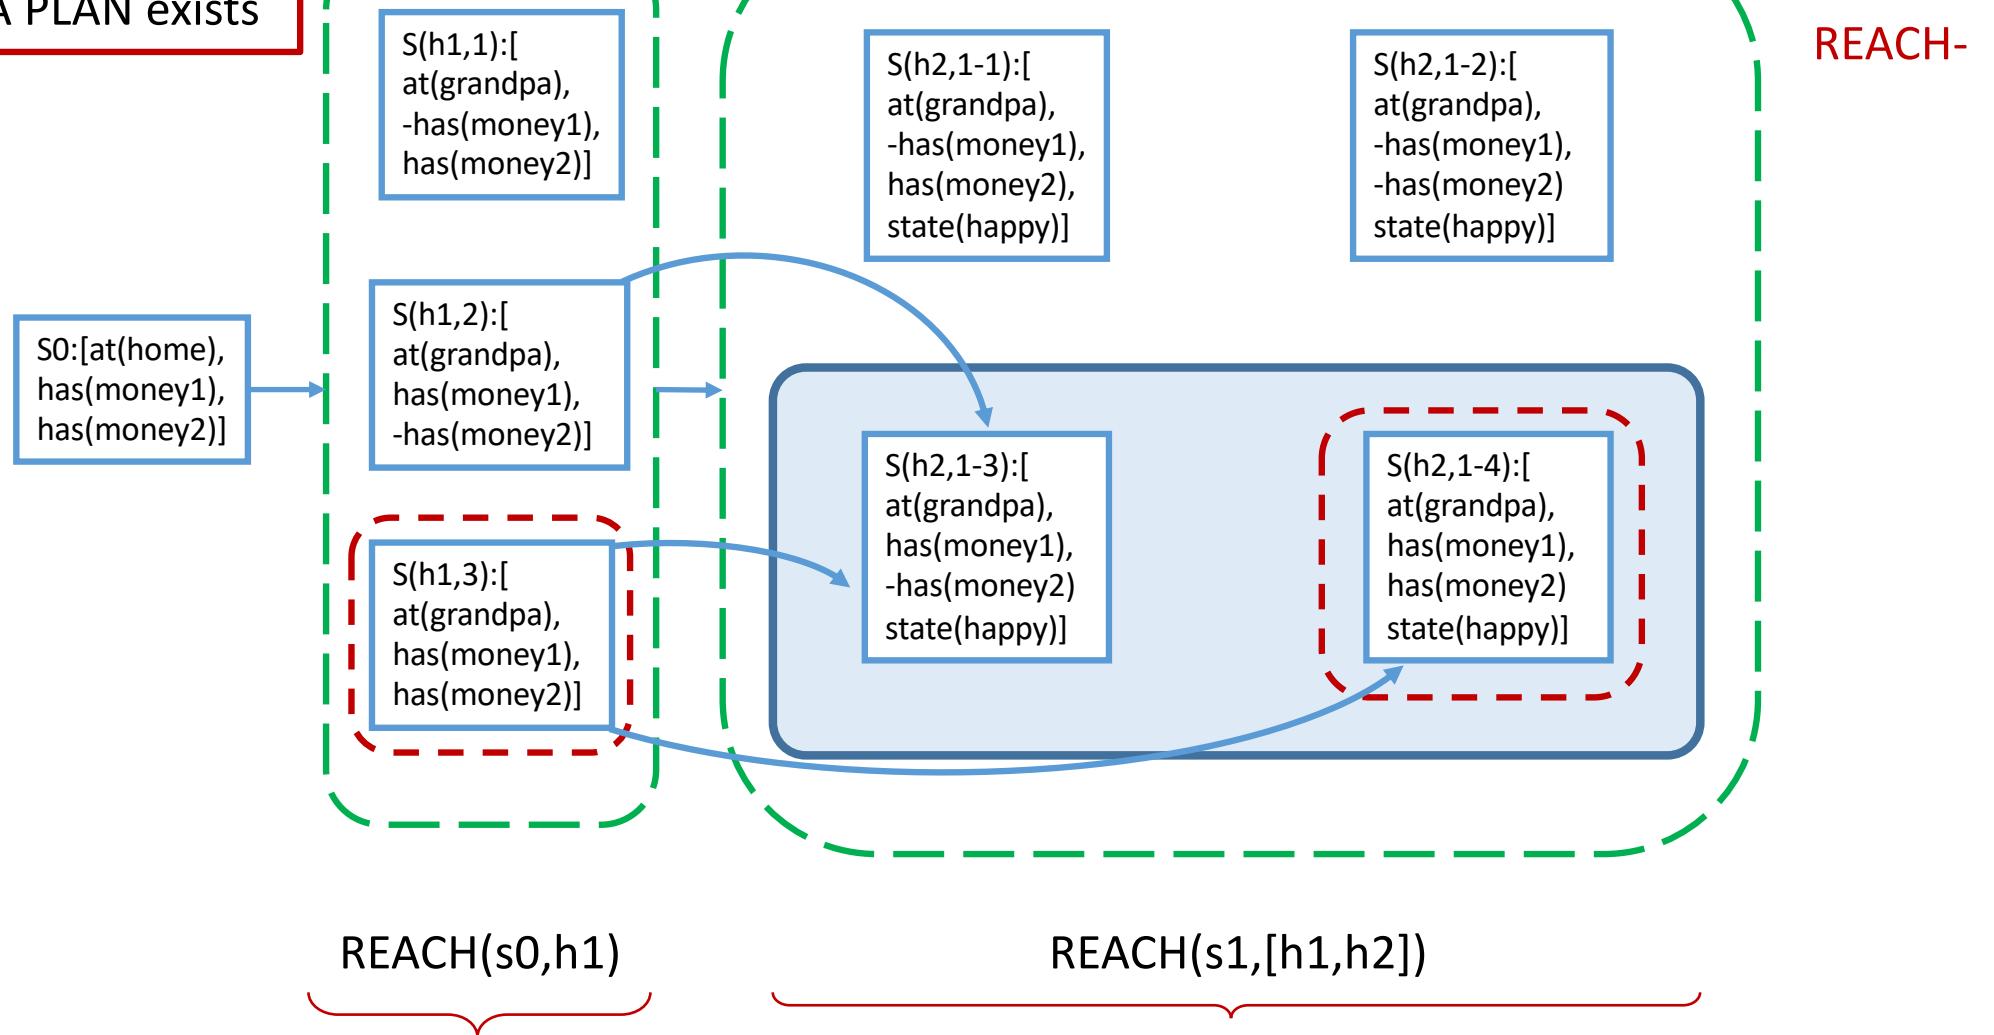
Visit Grandpa

$\text{REACH}(s,h)$

Initial state: $\text{has}(\text{money1}) \wedge \text{has}(\text{money2}) \wedge \text{at}(\text{home})$
 Goal state: $\text{has}(\text{money1}) \wedge \text{state}(\text{happy})$



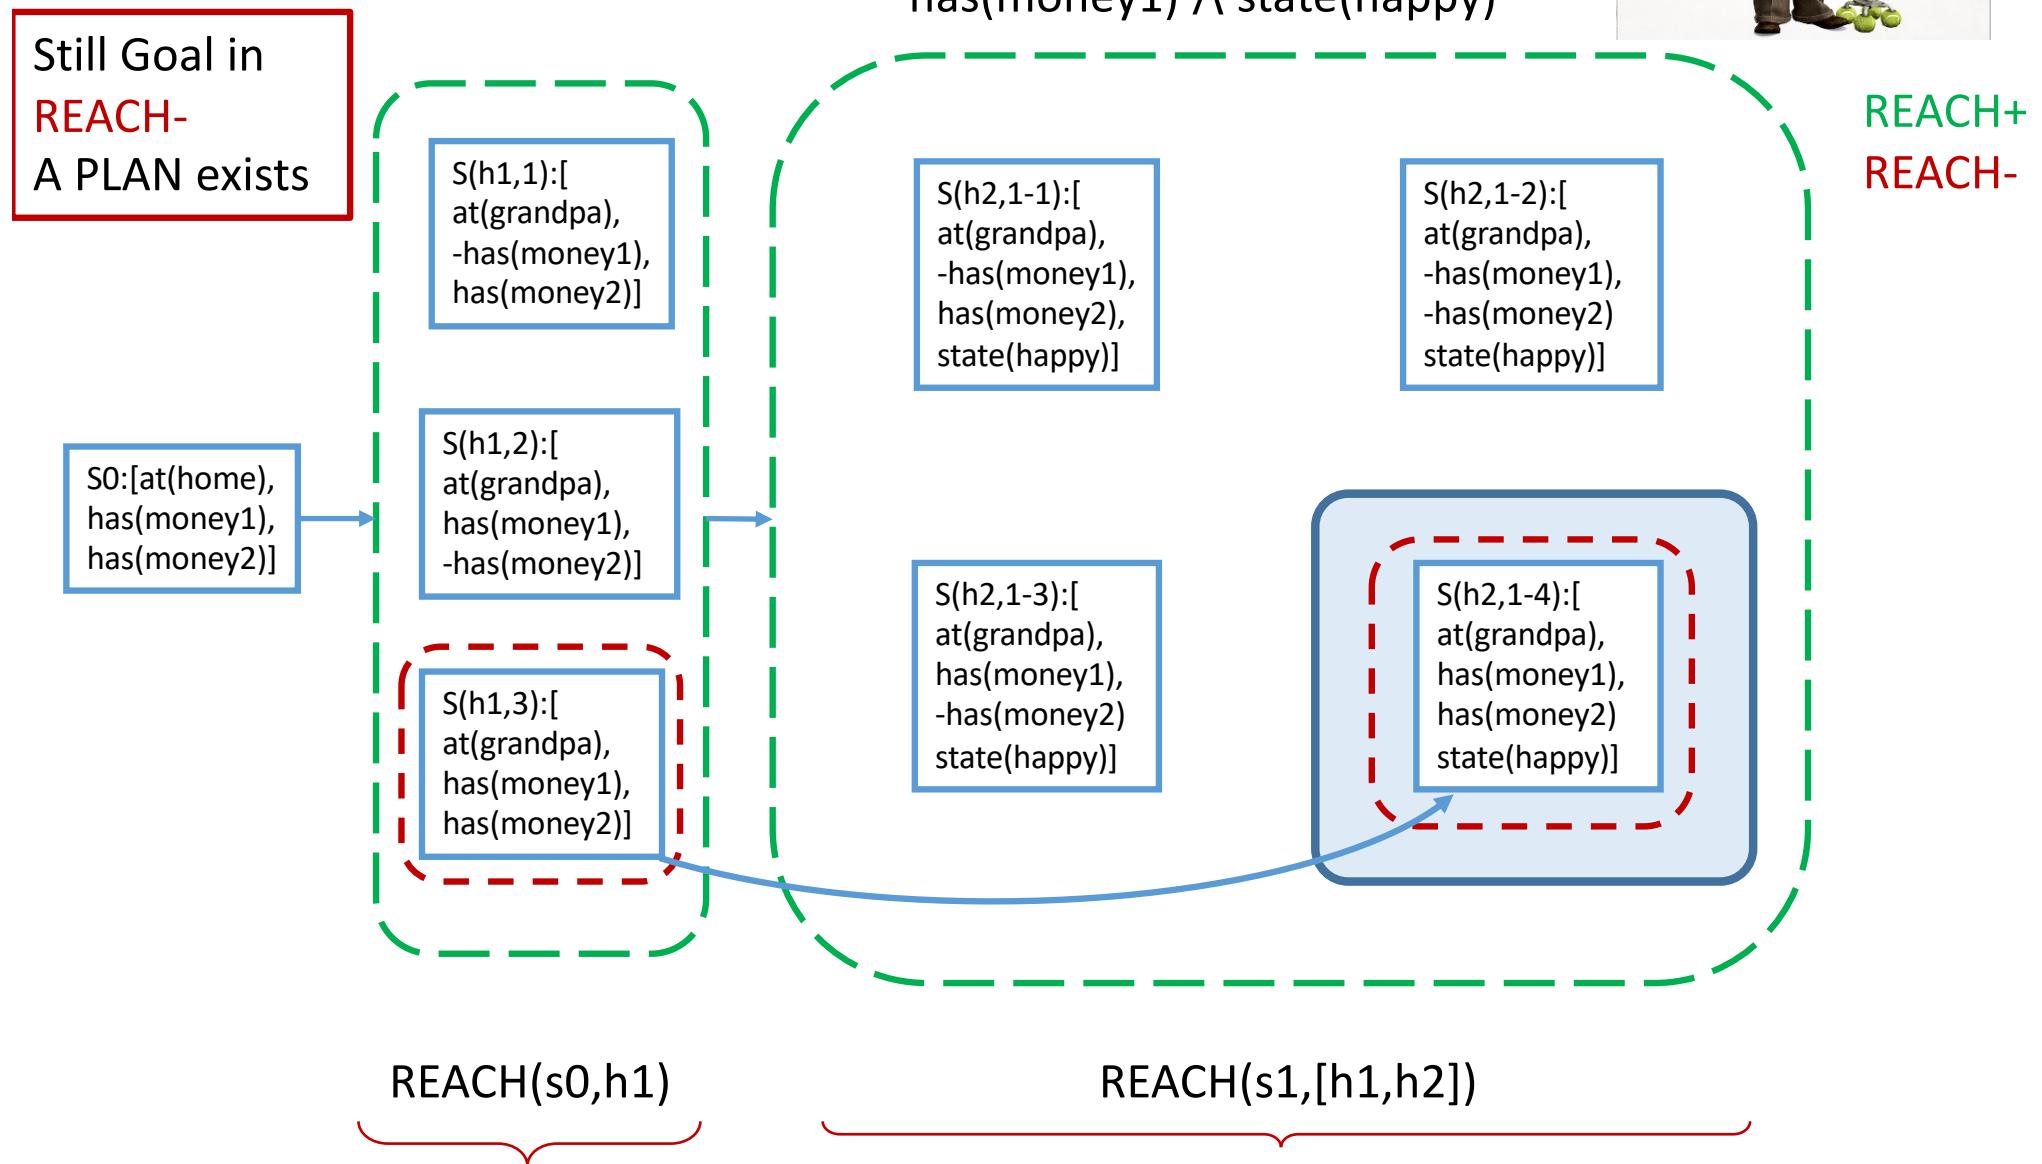
Still Goal in
 REACH-
 A PLAN exists



Visit Grandpa

$\text{REACH}(s,h)$

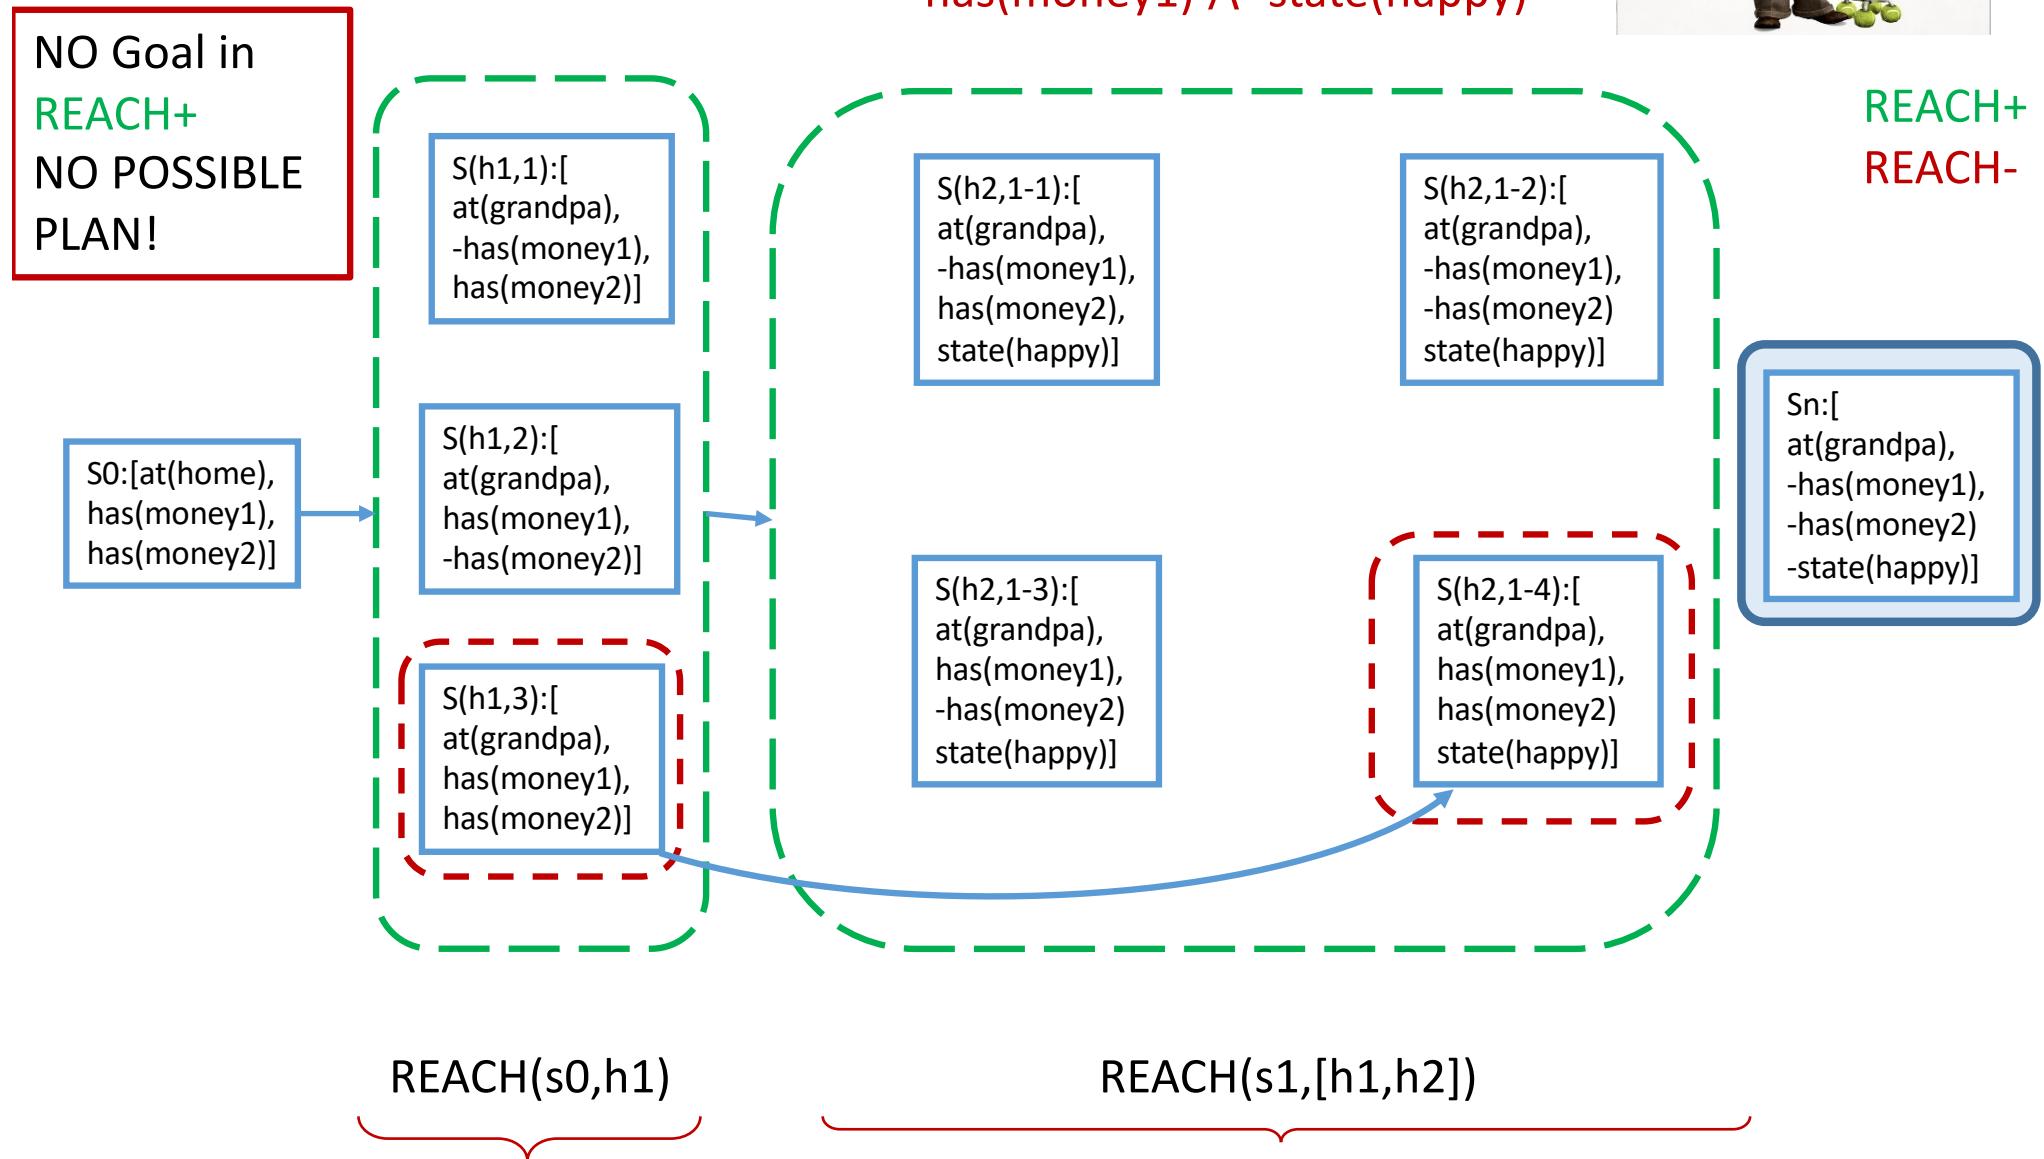
Initial state: $\text{has(money1)} \wedge \text{has(money2)} \wedge \text{at(home)}$
 Goal state: $\text{has(money2)} \wedge \text{has(money1)} \wedge \text{state(happy)}$



Visit Grandpa

$\text{REACH}(s,h)$

Initial state: $\text{has(money1)} \wedge \text{has(money2)} \wedge \text{at(home)}$
 Goal state: $\neg \text{has(money2)} \wedge \neg \text{has(money1)} \wedge \neg \text{state(happy)}$



Soccer Striker

Primitives:

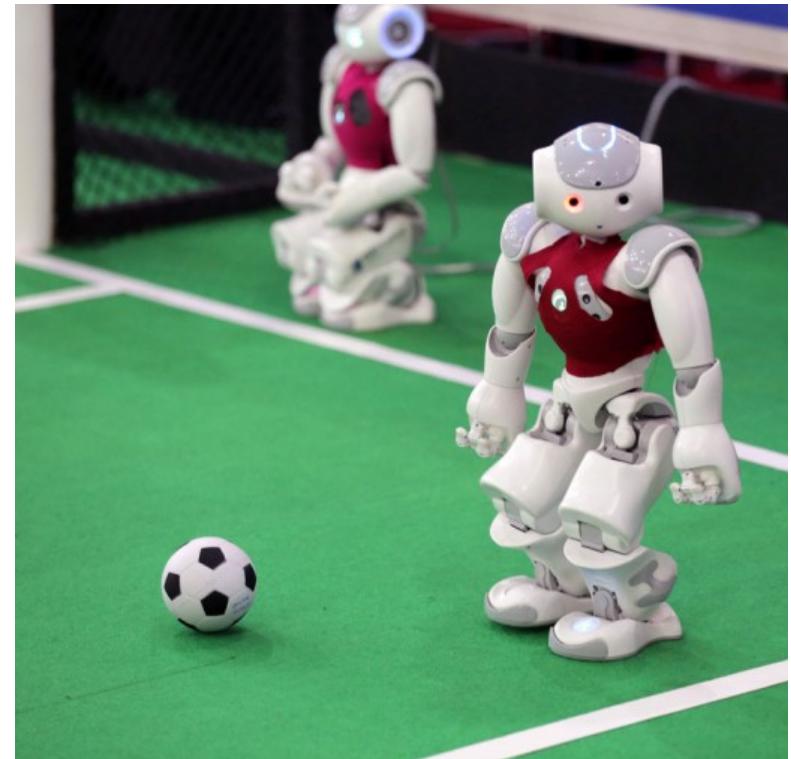
turn(left);
turnAround(left);
moveForward();

turn(right);
turnAround(right);
kick();

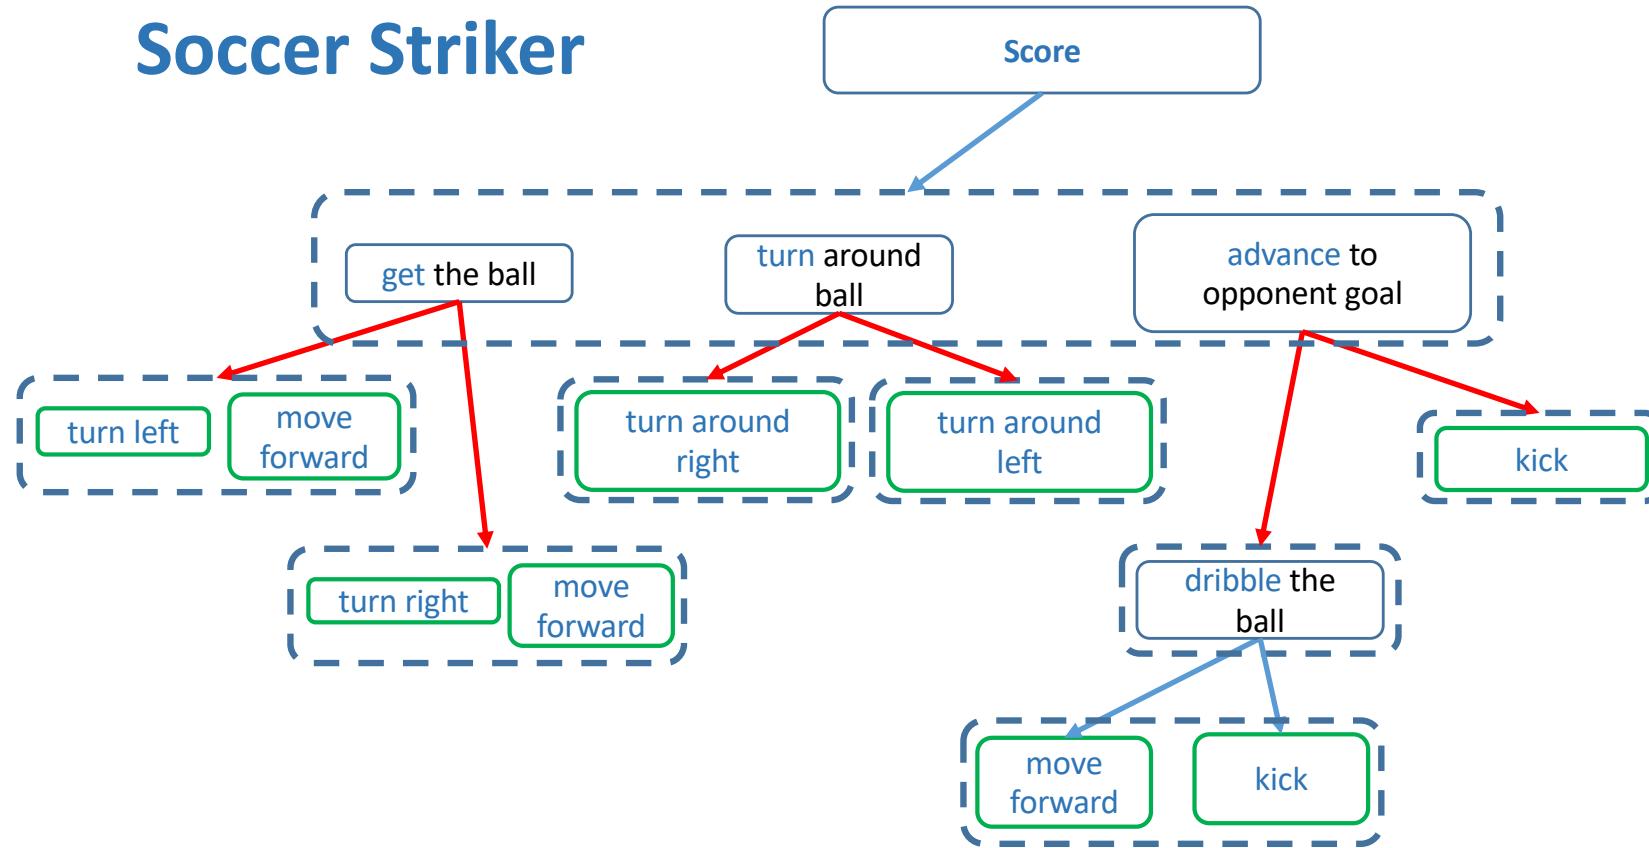
HLAs:

get the ball;
advance;

turn to goal;
dribble;



Soccer Striker



Soccer Striker

Primitives:

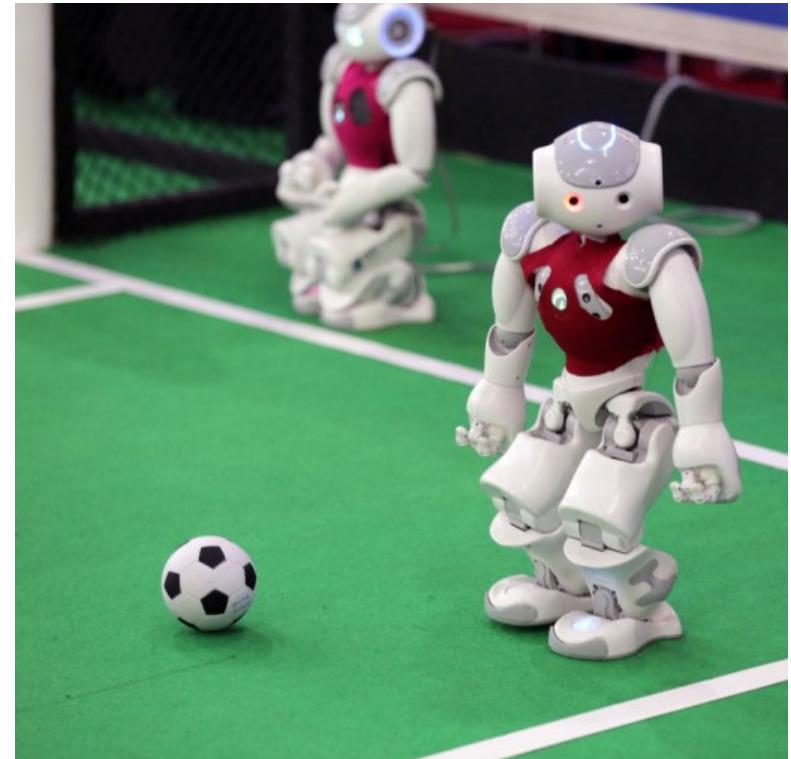
turn(left);
turnAround(left);
moveForward();
backKick();

turn(right);
turnAround(right);
kick();

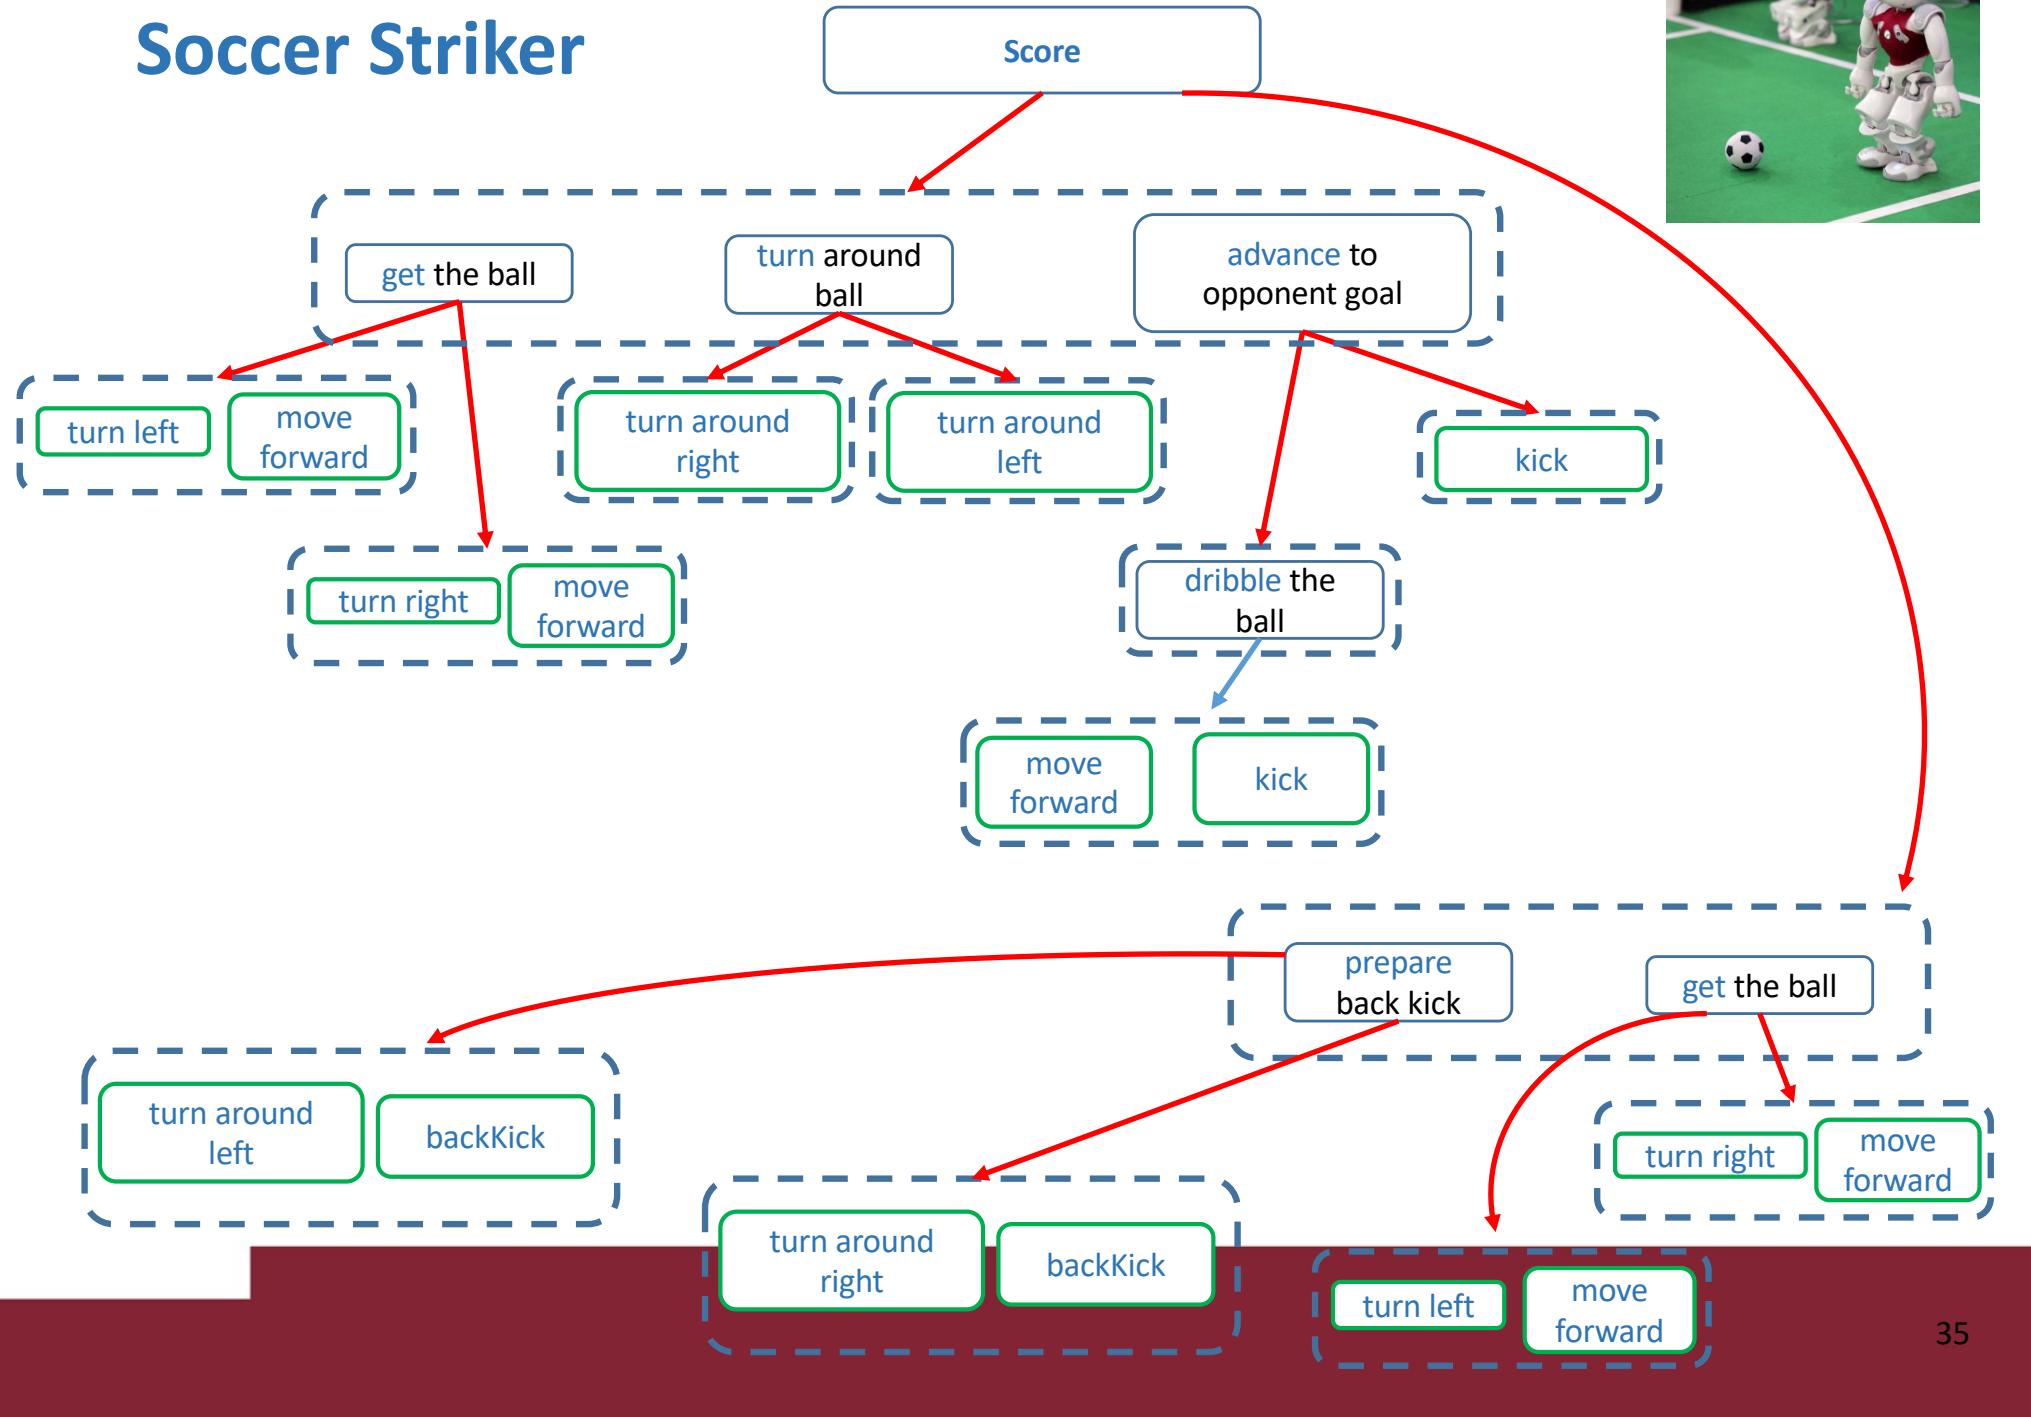
HLAs:

get the ball;
advance;
prepare back kick;

turn to goal;
dribble;



Soccer Striker



Recursive HTN Plans: example

Write the HLA `ClearStack(Block)` that takes all the blocks from the stack and puts them on the table.

Use it to create a plan for this problem:

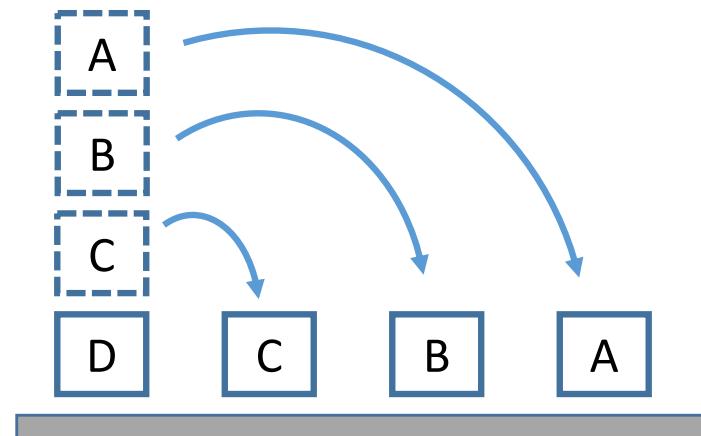
HLA

Refinement (`ClearStack(Block)`),
PREC: `on(Block, table)`
STEPS: [])

Refinement (`ClearStack(Block)`),
PREC: `-clear(Block) \wedge on(BlockAbove, Block)`
STEPS: [`ClearStack(BlockAbove)`,
`putOnTable(Block)`])

Init: `Clear(A) \wedge On(A,B) \wedge`
`On(B,C) \wedge On(C,D) \wedge`
`On(D,Table)`

Goal: `On(A,Table) \wedge On(B,Table)`
 `\wedge On(C,Table) \wedge On(D,Table)`



Action (`putOnTable(Block)`),
Prec: `clear(Block) \wedge on(Block, BlockBelow)`
Eff: `clear(BlockBelow) \wedge on(Block,table) \wedge -
on(Block, BlockBelow))`

A